



Real-World RAG: Eigene Daten & Dokumente mit semantischer Suche & LLMs erschließen

Sebastian Gingter
sebastian.gingter@thinktecture.com
Developer Consultant

Real-World RAG:

Eigene Daten & Dokumente mit semantischer Suche & LLMs erschließen

- Was Sie **ERWARTET**
 - Hintergrundwissen und Theorie zu RAG
 - Überblick über Semantische Suche
 - Probleme die auftreten können
 - Pragmatische Methoden für die Verwendung eigener Daten im RAG
 - Demos (Python)
- Was Sie **NICHT erwartet**
 - ChatGPT, CoPilot(s)
 - Grundlagen von ML
 - Deep Dives in LLMs, Vektor-Datenbanken, LangChain

Sebastian Gingter

Developer Consultant @ Thinktecture AG

- Generative AI in business settings
- Flexible and scalable backends
- All things .NET
- Pragmatic end-to-end architectures
- Developer productivity
- Software quality



Special Day

Generative AI für Business-Anwendungen

Thema	Sprecher	Datum, Uhrzeit
Large Language Models: Typische Use Cases & Patterns für Business-Anwendungen - in Action	Christian Weyer	DI, 17. September 2024, 10.45 bis 11.45
Real-World RAG: Eigene Daten & Dokumente mit semantischer Suche & LLMs erschließen	Sebastian Gingter	DI, 17. September 2024, 12.15 bis 13.15
Von 0 zu Smart: SPAs mit Generative AI aufwerten	Max Marschall	DI, 17. September 2024, 15.30 bis 16.30
Deep Dive in OpenAI Hosted Tools	Rainer Stropek	DI, 17. September 2024, 17.00 bis 18.00

Was Euch erwartet (und was nicht):

- Ein bisschen Hintergrund-Info & Theorie
- Überblick über das Themengebiet Semantische Suche
- Probleme und mögliche Strategien
- Pragmatische Ansätze für die eigenen Daten
- Kein C#, sondern Python 🤖
- Kein Deep-Dive in
 - LLMs
 - Vektor-Datenbanken
 - LangChain

Agenda

- Short Introduction to RAG
- Embeddings (and a bit of theory 🤖)
- Vector-Databases
- Indexing
- Retrieval
- Not good enough? – Indexing II
 - HyDE & alternative indexing methods
- Conclusion

Introduction

Introduction

Embeddings

Vector-DBs


Indexing

Retrieval

Indexing II




RAG

Use-case: Talk to my internal data

 Ask or search...

Darn, I missed my last connecting train and won't come home today. What can I do?

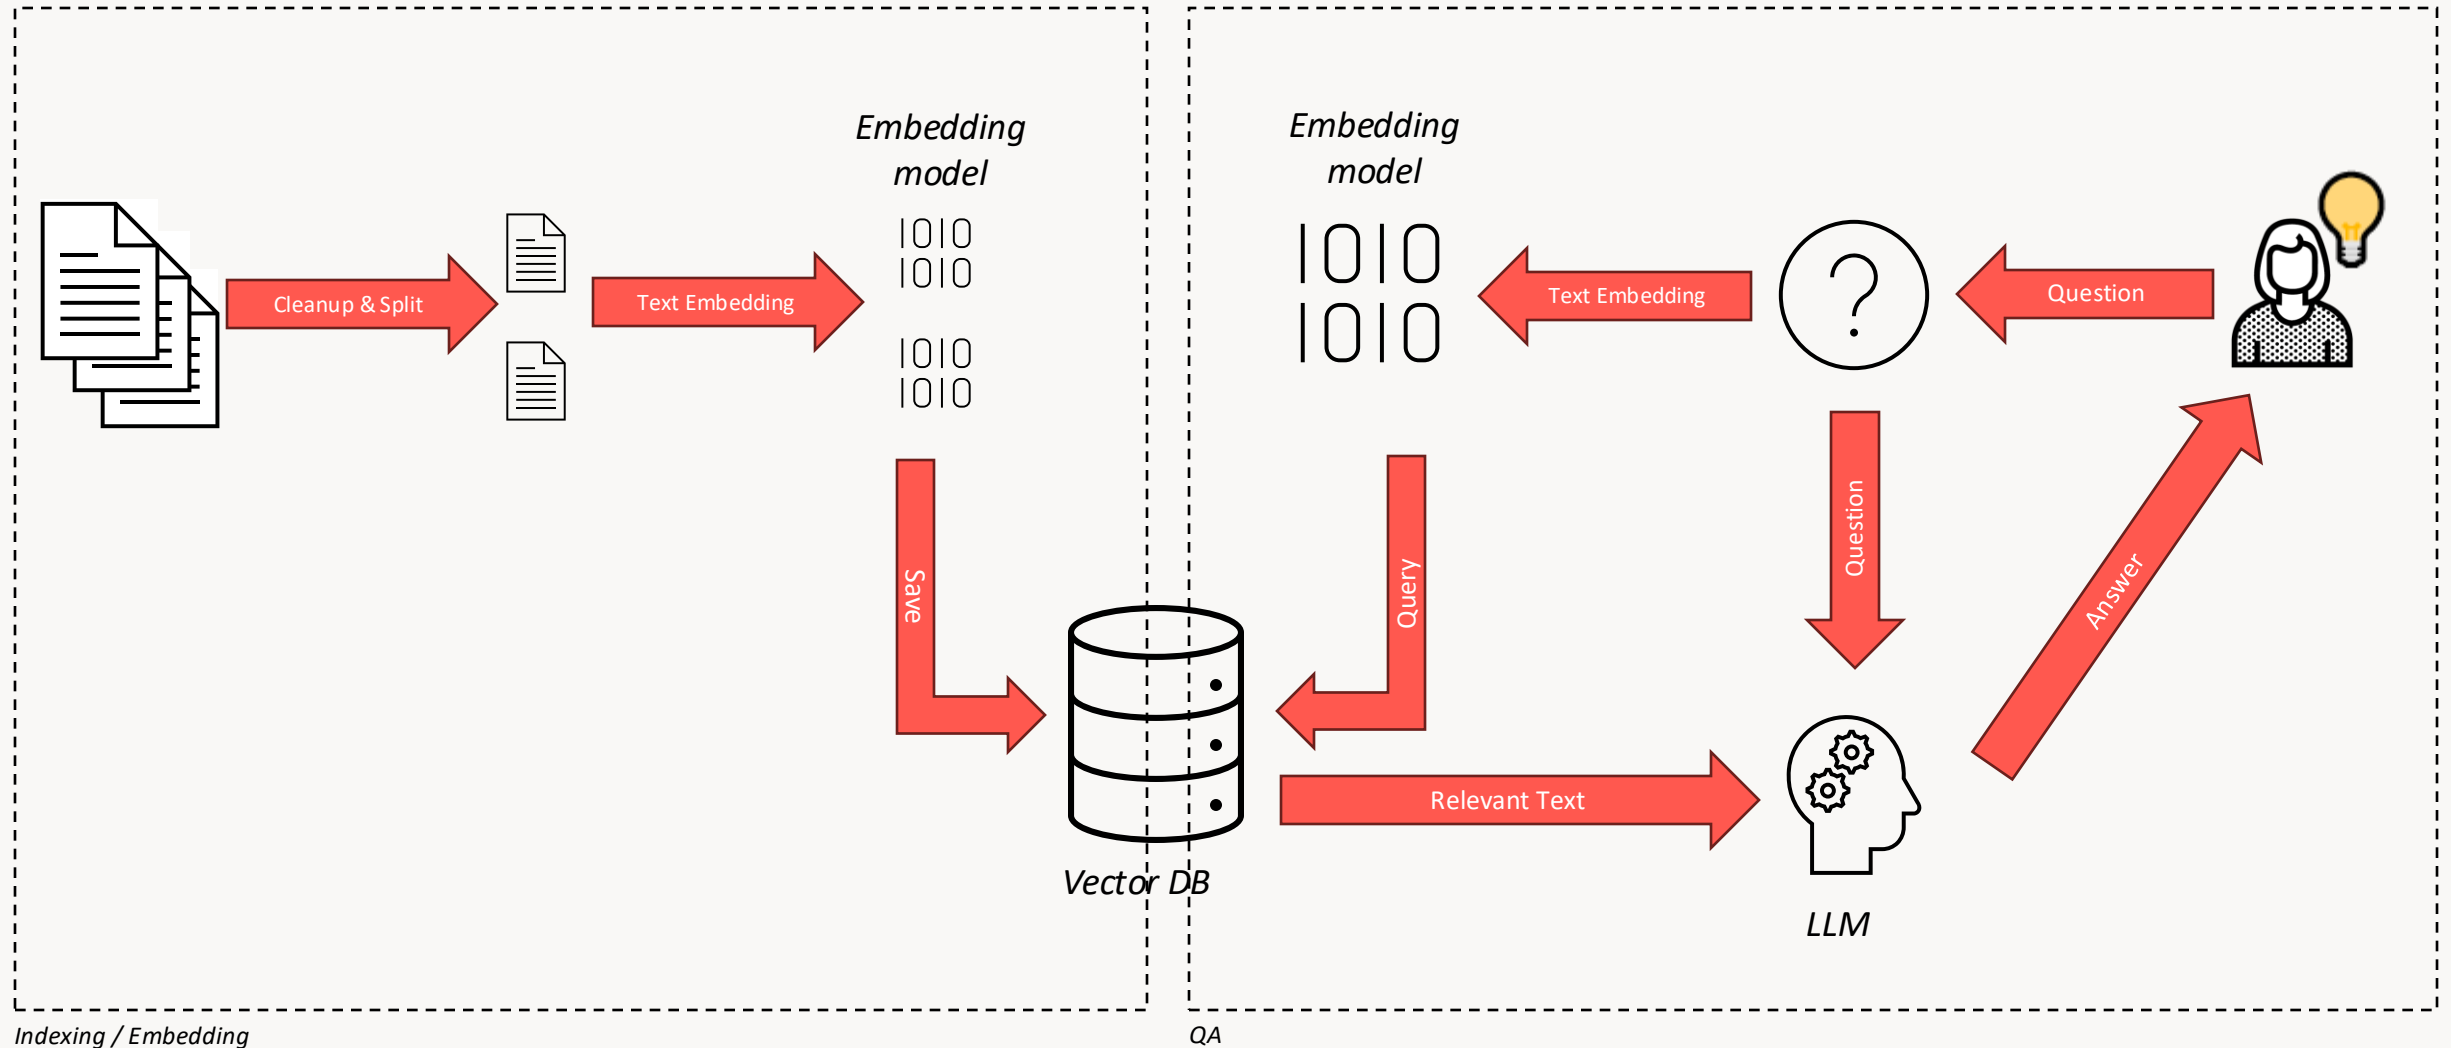
If your flight is delayed and you miss the last train home or to the client, you can look for a nearby hotel or check for a taxi, making sure to ask for a fixed price before starting the journey. Expenses up to 300-400 EUR are generally not a problem for such situations.



Answer based on 1 sources >

Retrieval-augmented generation (RAG)

Indexing & (Semantic) search



Semantic Search

- Classic search: lexical
 - Compares words, parts of words and variants
 - Classic SQL: WHERE 'content' LIKE '%searchterm%'
 - We can search only for things where we know that its somewhere in the text
- New: Semantic search
 - Compares for the same contextual meaning
 - "Das Rudel rollt das runde Gerät auf dem Rasen herum"
 - "The pack enjoys rolling a round thing on the green grass"
 - "Die Hunde spielen auf der Wiese mit dem Ball"
 - "The dogs play with the ball on the meadow"

Semantic Search

- How to grasp “semantics”?
- Computers only calculate on numbers
 - Computing is “applied mathematics”
- AI also only calculates on numbers
- We need a numeric representation of meaning
 - ➔ “Embeddings”

Embeddings

Introduction

Embeddings

Vector-DBs

Indexing

Retrieval

Indexing II

RAG

Embedding (math.)

- Topologic: Value of a high dimensional space is “embedded” into a lower dimensional space
- Natural / human language is very complex (high dimensional)
 - Task: Map high complexity to lower complexity / dimensions
- Injective function
- Similar to hash, or a lossy compression

Embeddings


- Embedding model (specialized ML model) converting text into a numeric representation of its meaning
- Representation is a vector in an n-dimensional space
 - n floating point values
 - OpenAI
 - “text-embedding-ada-002” uses 1536 dimensions
 - “text-embedding-3-small” 512 and 1536
 - “text-embedding-3-large” 256, 1024 and 3072
 - Huggingface models have a very wide range of dimensions

Embeddings

- Embedding models are unique
- Each dimension has a different meaning, individual to the model
- vectors from different models are incompatible with each other
- Some embedding models are multi-language, but not all
- In an LLM, also the first step is to embed the input into a lower dimensional space

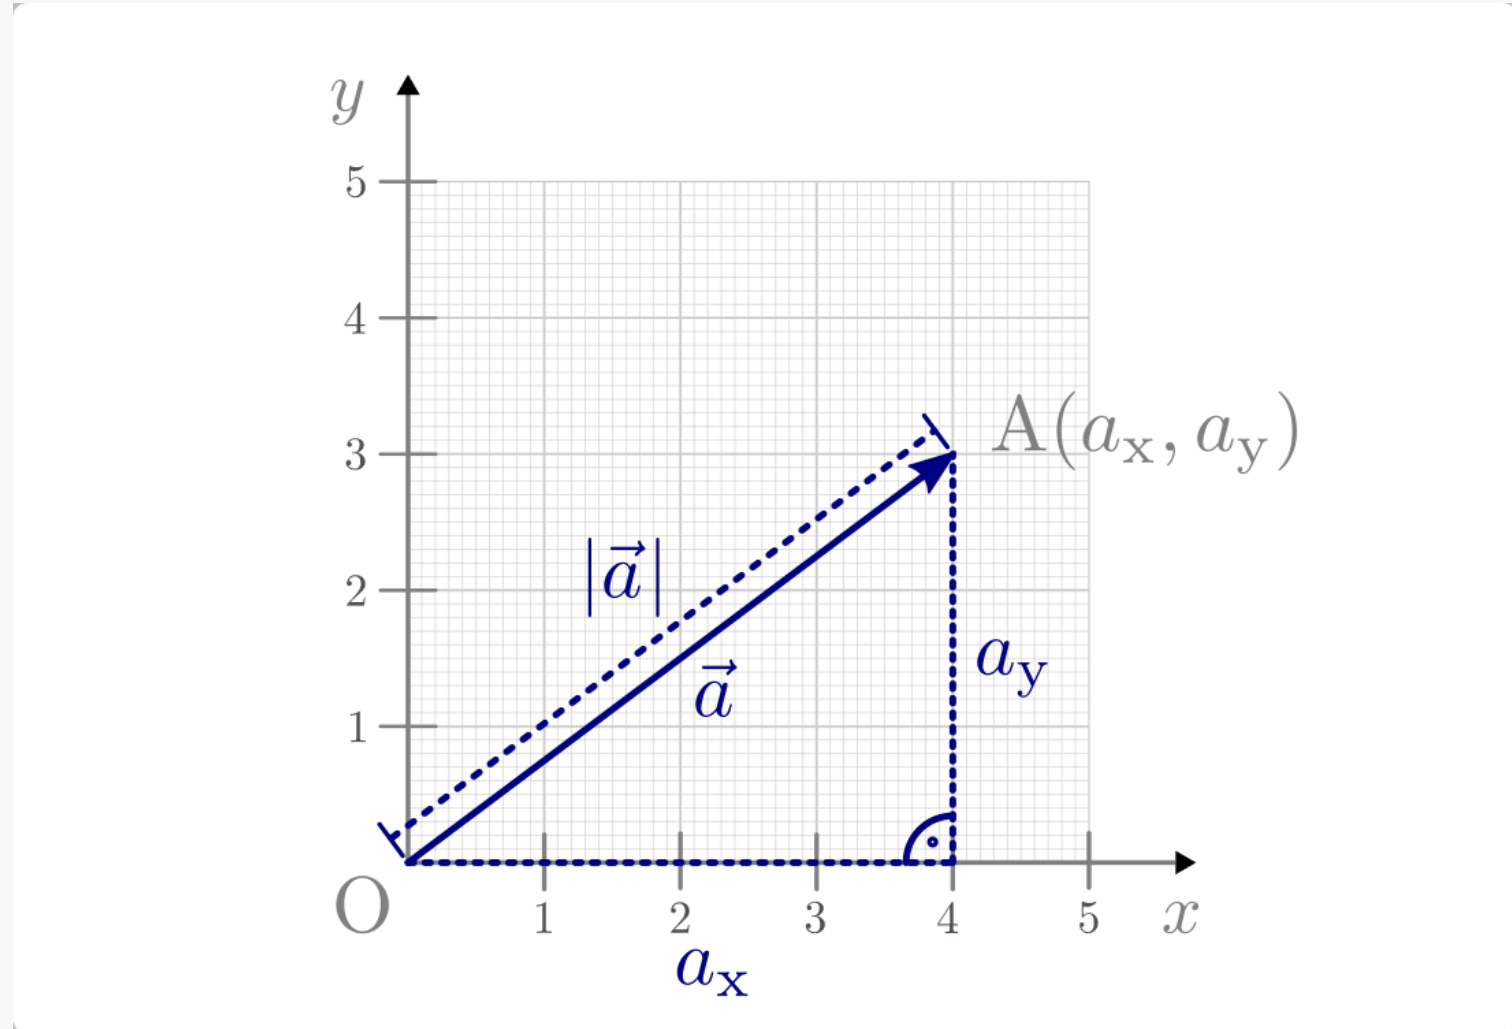
What is a vector?

- Mathematical quantity with a direction and length
- $\vec{a} = \begin{pmatrix} a_x \\ a_y \end{pmatrix}$




Vectors in 2D

$$\vec{a} = \begin{pmatrix} a_x \\ a_y \end{pmatrix}$$



Vectors in 3D

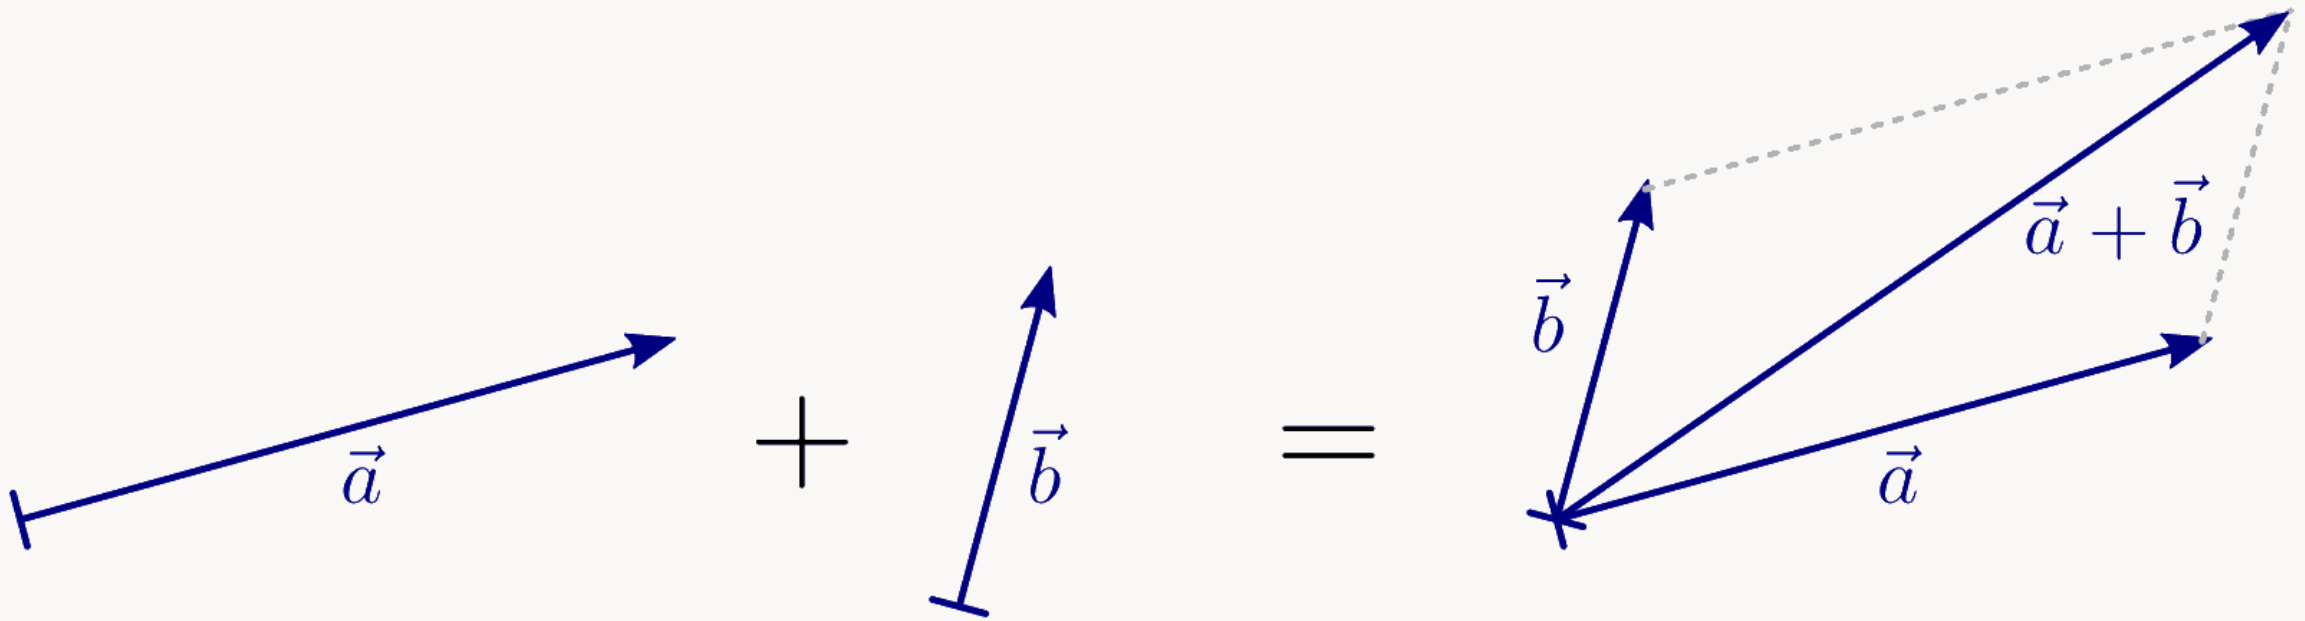
$$\vec{a} = \begin{pmatrix} a_x \\ a_y \\ a_z \end{pmatrix}$$



Vectors in multidimensional space

$$\vec{a} = \begin{pmatrix} a_u \\ a_v \\ a_w \\ a_x \\ a_y \\ a_z \end{pmatrix}$$

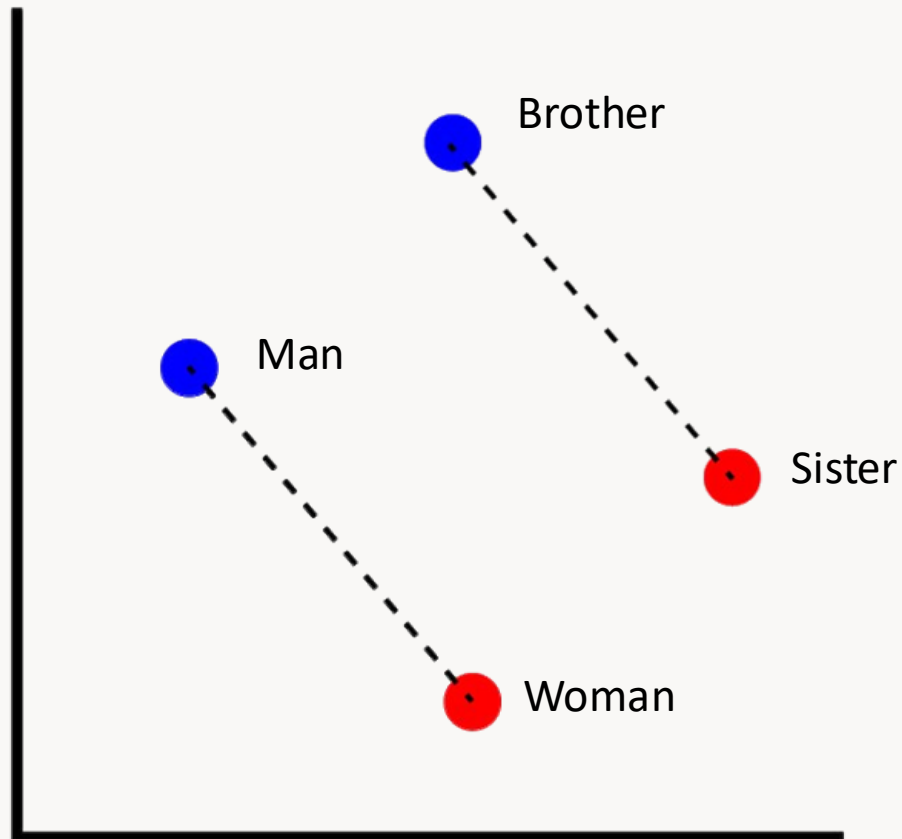
Calculation with vectors



Word2Vec

Mikolov et al., Google, 2013


$$\textit{Brother} - \textit{Man} + \textit{Woman} \approx \textit{Sister}$$



Embedding-Model


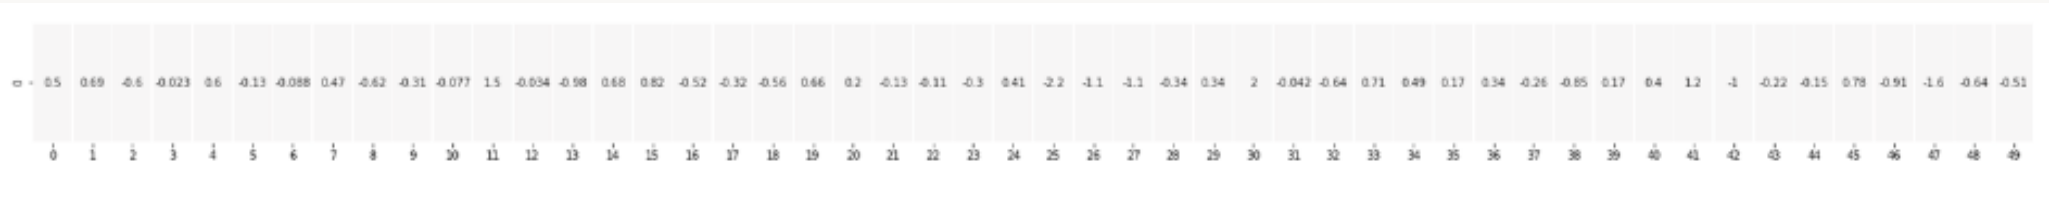
- Task: Create a vector from an input
 - Extract meaning / semantics
- Embedding models usually are very shallow & fast
Word2Vec is only two layers
- Similar to the first step of an LLM
 - Convert text to values for input layer
- This comparison is very simplified, but one could say:
 - The embedding model 'maps' the meaning into the model's 'brain'

Embedding-Model




Embedding-Model

[0.50451, 0.68607, -0.59517, -0.022801, 0.60046, -0.13498, -0.08813, 0.47377, -0.61798, -0.31012, -0.076666, 1.493, -0.034189, -0.98173, 0.68229, 0.81722, -0.51874, -0.31503, -0.55809, 0.66421, 0.1961, -0.13495, -0.11476, -0.30344, 0.41177, -2.223, -1.0756, -1.0783, -0.34354, 0.33505, 1.9927, -0.04234, -0.64319, 0.71125, 0.49159, 0.16754, 0.34344, -0.25663, -0.8523, 0.1661, 0.40102, 1.1685, -1.0137, -0.21585, -0.15155, 0.78321, -0.91241, -1.6106, -0.64426, -0.51042]



Embedding-Model



Recap Embeddings

- Embedding model: “Analog to digital converter for text”
- Embeds the high-dimensional natural language meaning into a lower dimensional-space (the model’s ‘brain’)
- No magic, just applied mathematics
- Math. representation: Vector of n dimensions
- Technical representation: array of floating point numbers

DEMO

Embeddings

Sentence Transformers, local embedding model

Vector-Databases

Introduction

Embeddings

Vector-DBs

Indexing

Retrieval

Indexing II

RAG

Vector-Databases

- Mostly document-based
- Index: Embedding (vector)
- Document (content)
- Metadata
- Query functionalities

Vector-Databases

- Pinecone
- Milvus
- Chroma
- Weaviate
- Deep Lakee
- Qdrant
- Elasticsearch
- Vespa
- Vald
- ScaNN
- Pgvector
(PostgreSQL Extension)
- Faiss
- ...
- ... (probably) coming to a relational database near you soon(ish)
SQL Server Example: <https://learn.microsoft.com/en-us/samples/azure-samples/azure-sql-db-openai/azure-sql-db-openai/>

Vector-Databases

- (Search-)Algorithms

- Cosine Similarity $S_{C(a,b)} = \frac{a \cdot b}{\|a\| \times \|b\|}$
- Manhattan Distance (L1 norm, taxicab)
- Euclidean Distance (L2 norm)
- Minkowski Distance (~ generalization of L1 and L2 norms)
- L_∞ (L-Infinity), Chebyshev Distance
- Jaccard index / similarity coefficient (Tanimoto index)
- Nearest Neighbour
- Bregman divergence
- ...

DEMO

Vector database

LangChain, Chroma, local embedding model

Indexing

Introduction

Embeddings

Vector-DBs

Indexing

Retrieval

Indexing II

RAG









Indexing

- Loading
- Clean-up
- Splitting
- Embedding
- Storing

Loading

- Import documents from different sources, in different formats
- LangChain has very strong support for loading data
- Support for cleanup
- Support for splitting

Document loaders

 mhtml MHTML is a is used both for emails but also for archived webpag...	 Microsoft Excel The UnstructuredExcelLoader is used to load Microsoft Excel files.
 Microsoft OneDrive Microsoft OneDrive (formerly	 Microsoft OneNote This notebook covers how to load documents from OneNote.
 Microsoft PowerPoint [Microsoft	 Microsoft SharePoint Microsoft SharePoint is a
 Microsoft Word Microsoft Word	 Modern Treasury Modern Treasury simplifies complex


Clean-up

- HTML Tags
- Formatting information
- Normalization
 - lowercasing
 - stemming, lemmatization
 - remove punctuation & stop words
- Enrichment
 - tagging
 - keywords, categories
 - metadata



Splitting (Text Segmentation)

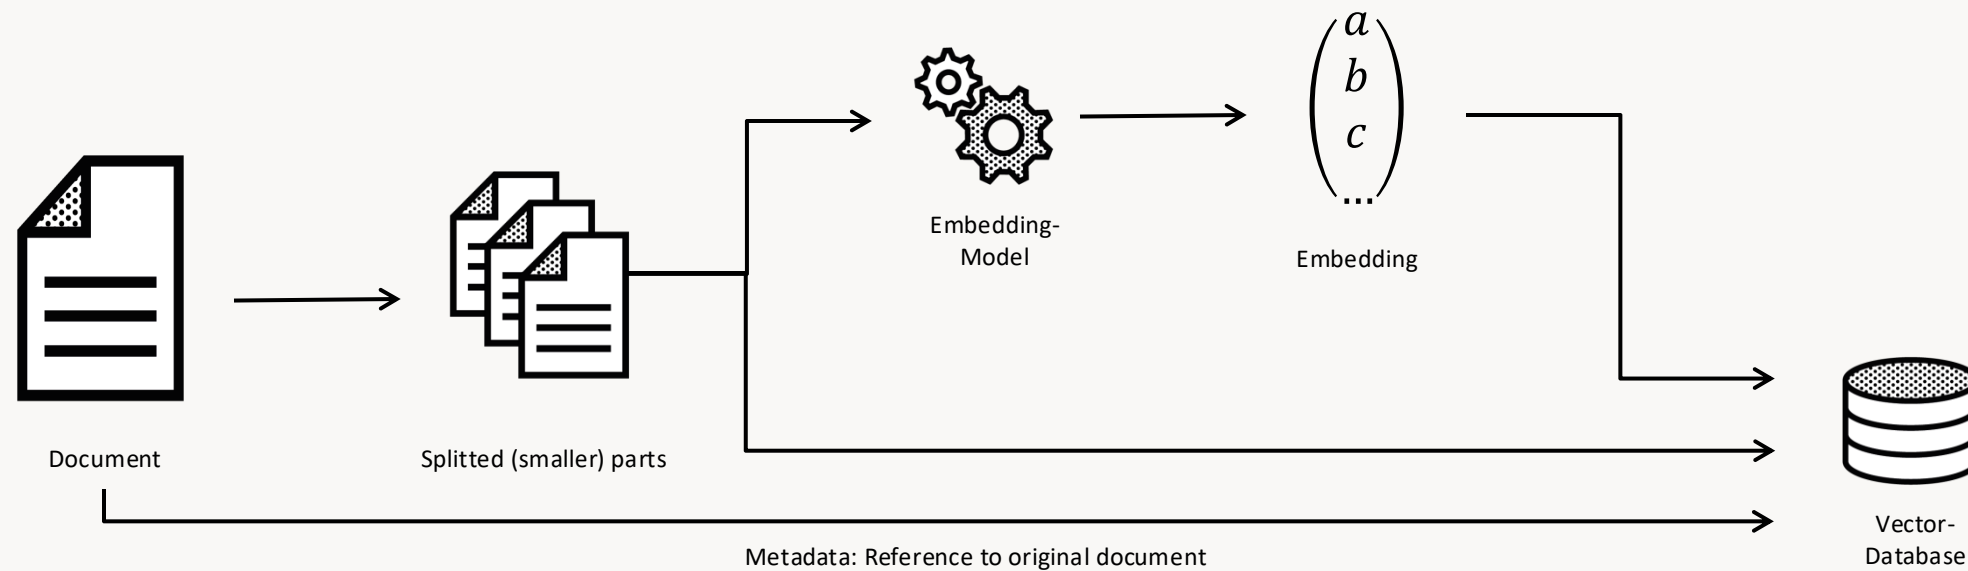
- Document is too large / too much content / not concise enough



- by size (text length)
- by character (`\n\n`)
- by paragraph, sentence, words (until small enough)
- by size (tokens)
- overlapping chunks (token-wise)

Vector-Databases

- Indexing



Retrieval (Search)

Introduction

Embeddings

Vector-DBs


Indexing

Retrieval

Indexing II

RAG

Retrieval



DEMO

Store and retrieval

LangChain, Chroma, local embedding model, OpenAI GPT

Indexing II

Not good enough?

Introduction

Embeddings

Vector-DBs


Indexing

Retrieval

Indexing II

RAG

Not good enough?

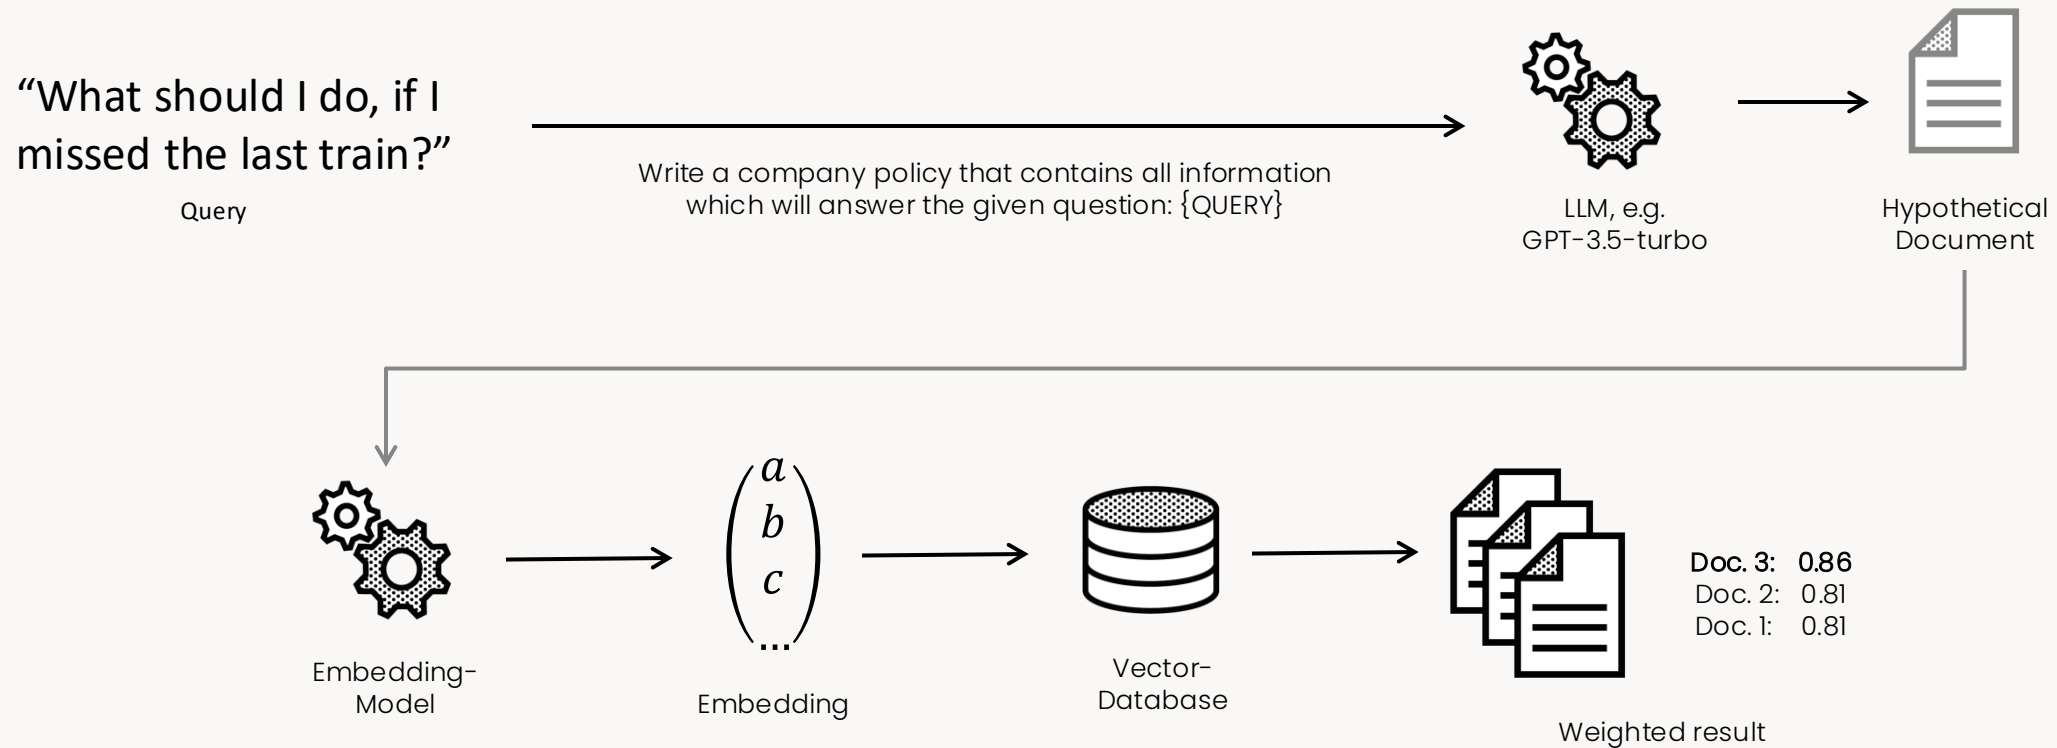


Not good enough?

- Semantic search still only uses your index
- It's just as good as your embeddings
 - All chunks need to be
- Sh*t in, sh*t out

HyDE (Hypothetical Document Embeddings)

- Search for a hypothetical Document

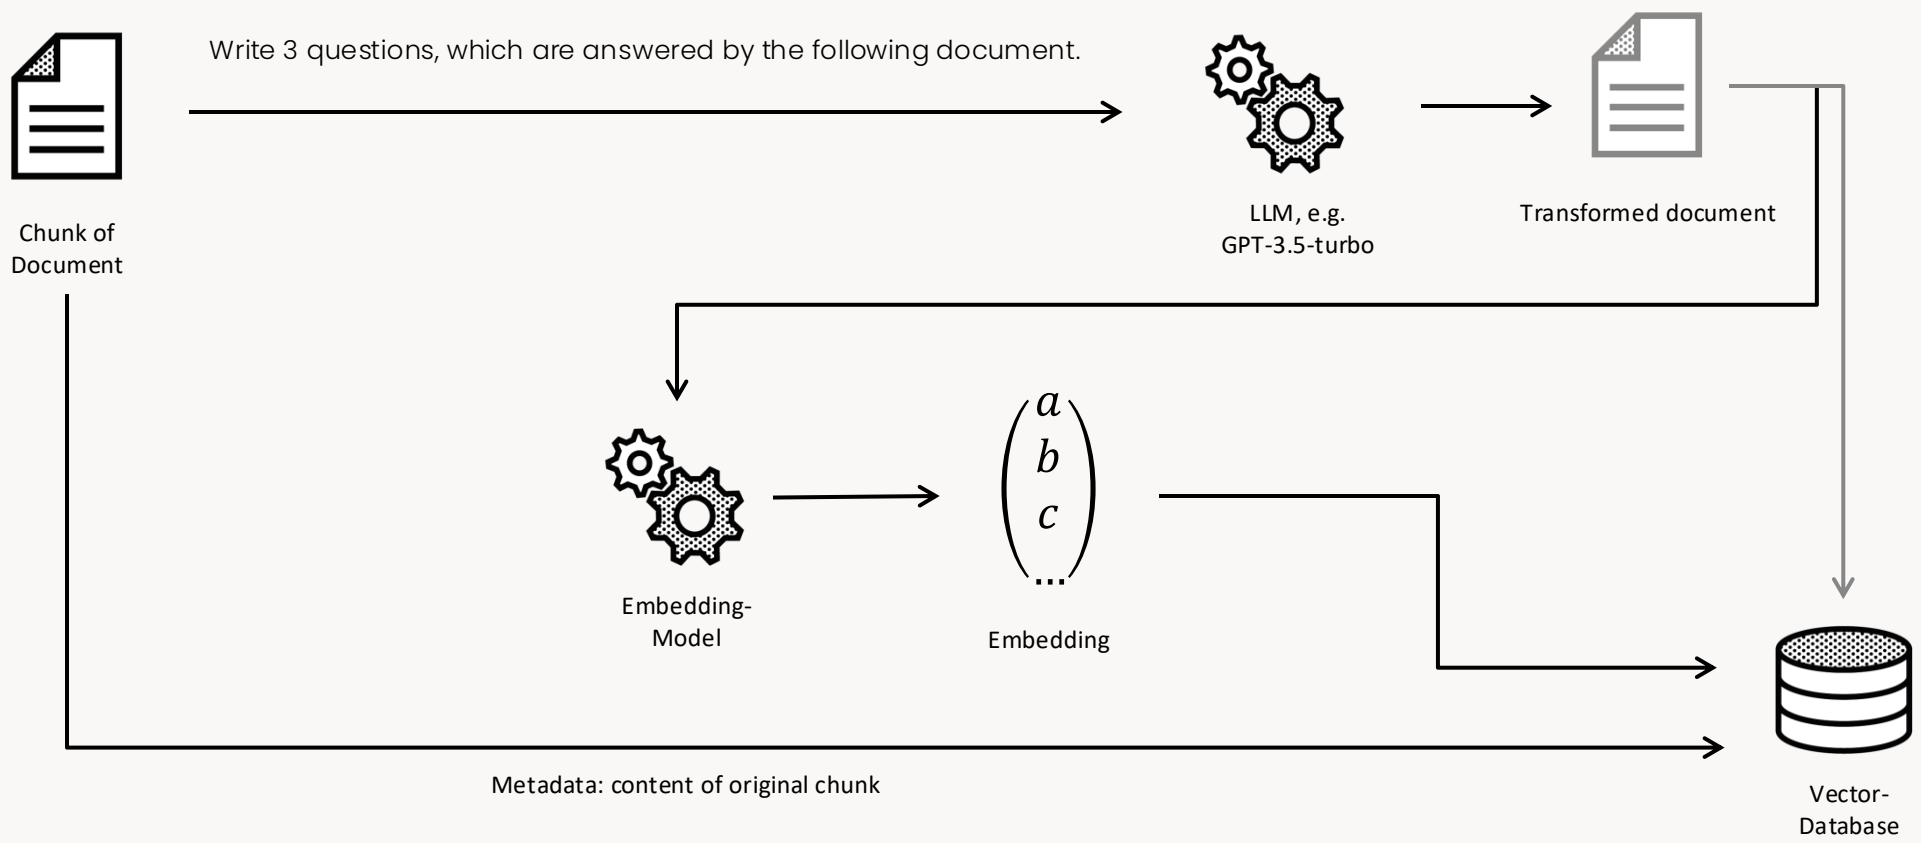


What else?

- Downside of HyDE:
 - Each request needs to be transformed through an LLM (slow & expensive)
 - A lot of requests will probably be very similar to each other
 - Each time a different hyp. document is generated, even for an extremely similar request
 - Leads to very different results each time
- Idea: Alternative indexing
 - Transform the document, not the query

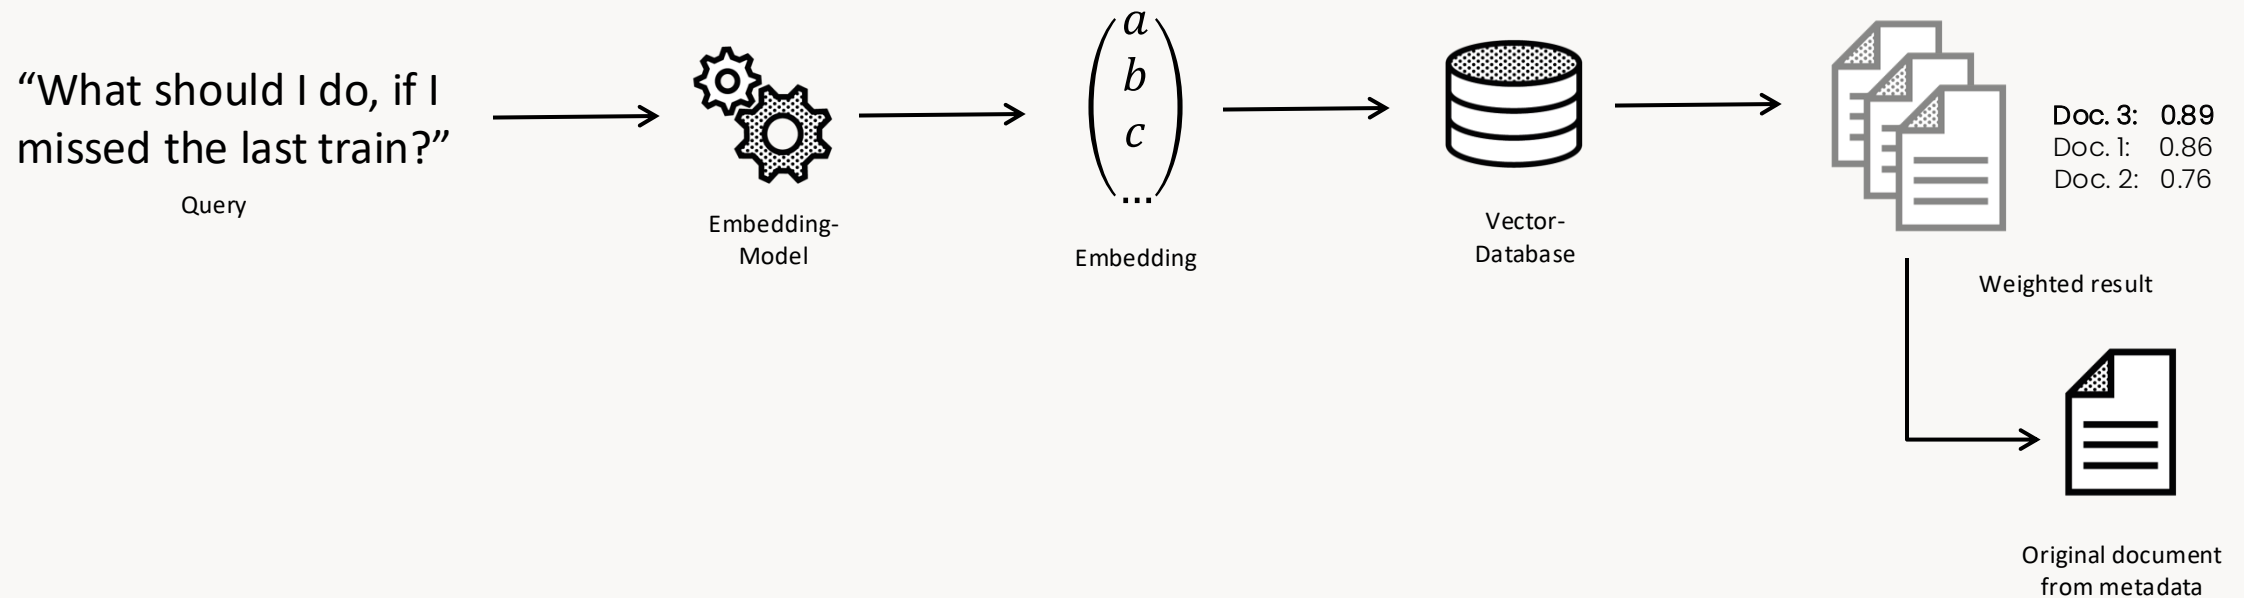
Alternative Indexing

HyQE: Hypothetical Question Embedding



Alternative Indexing

- Retrieval



DEMO

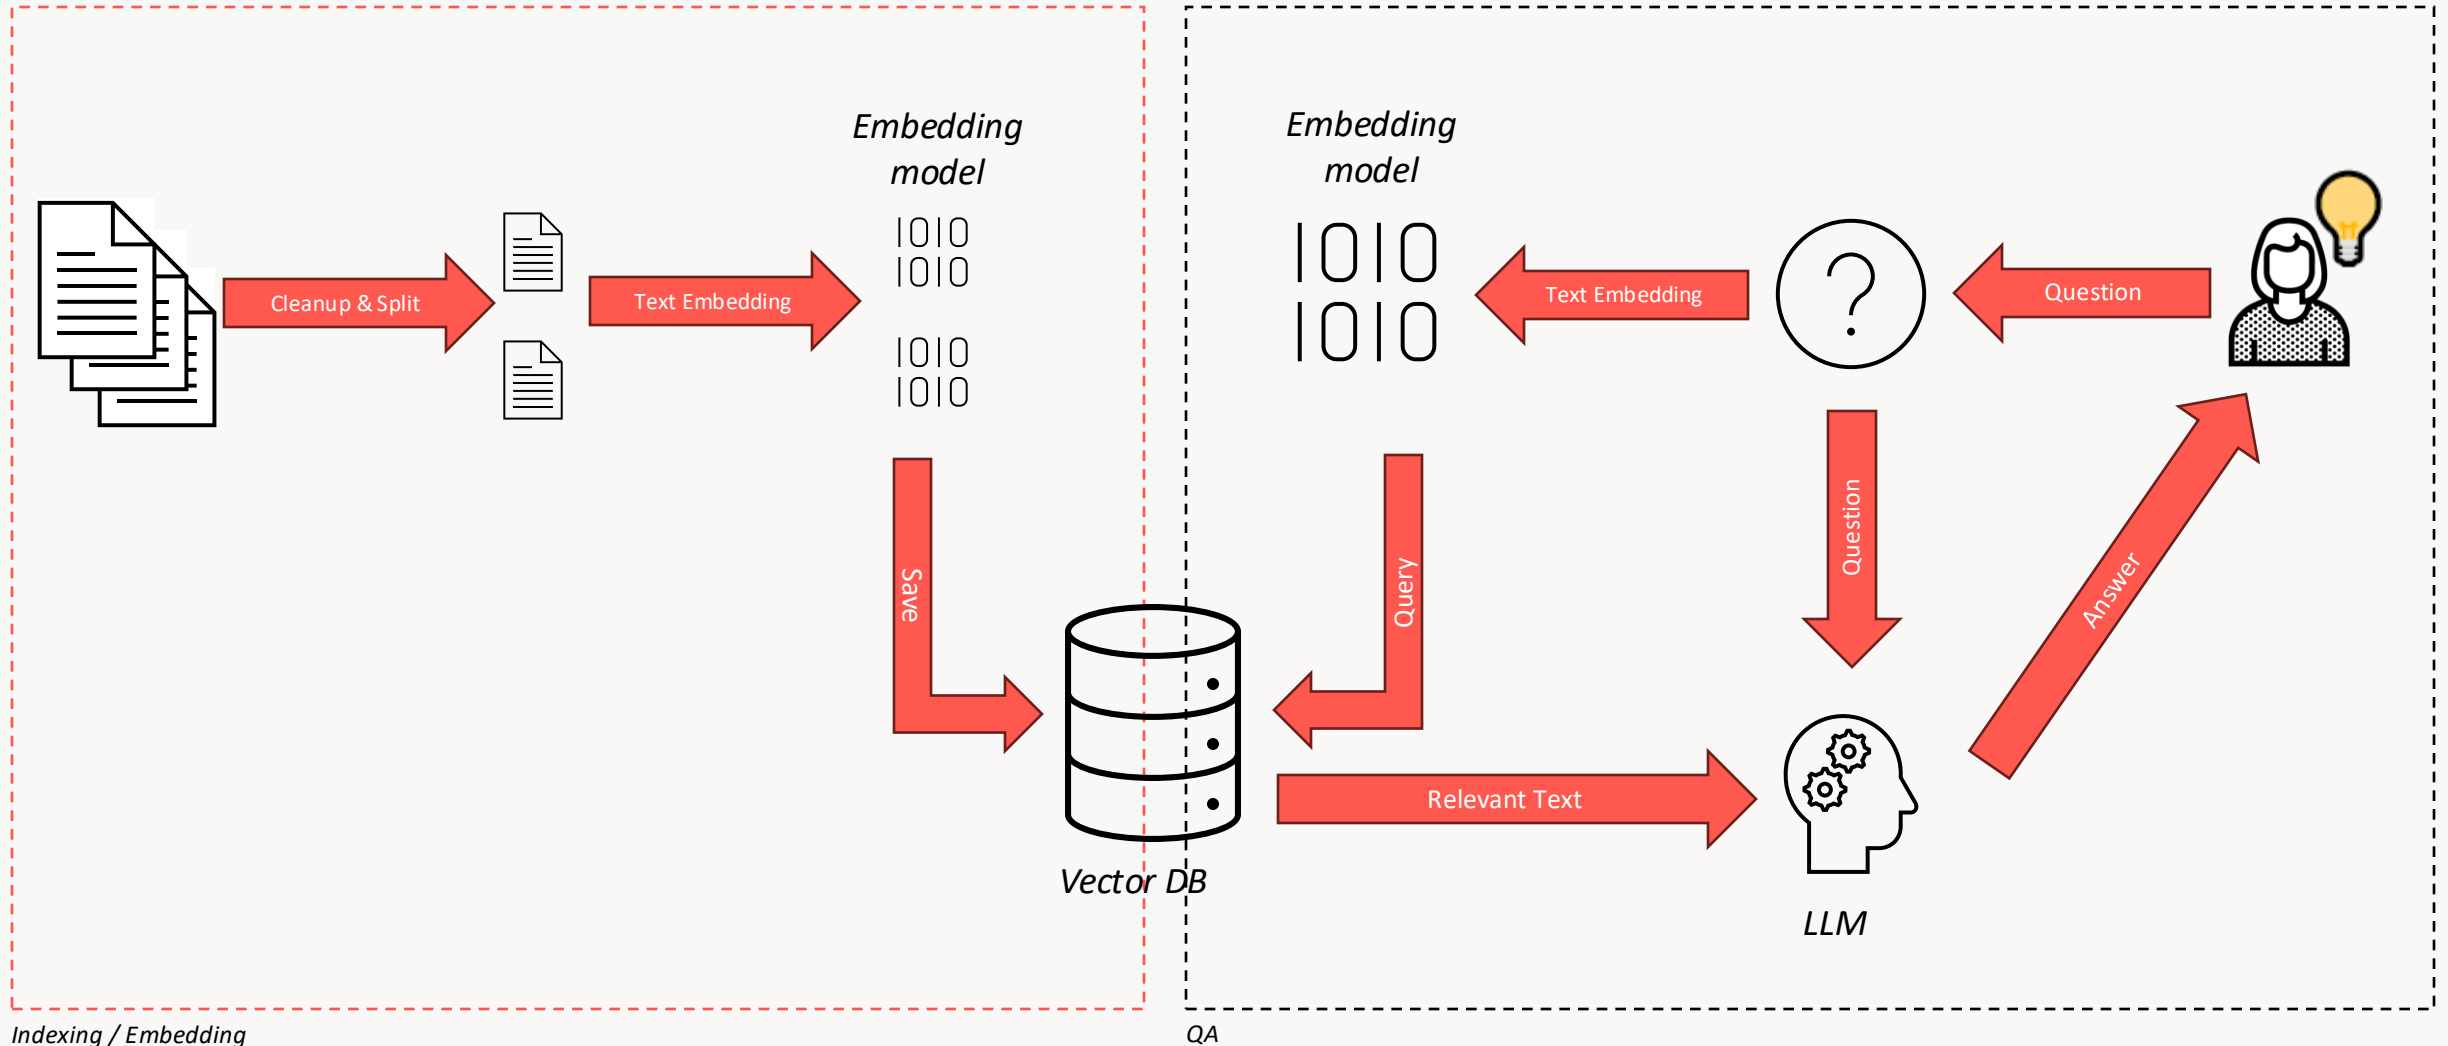
Compare embeddings

LangChain, Qdrant, OpenAI GPT

Conclusion

Retrieval-augmented generation (RAG)

Indexing & (Semantic) search

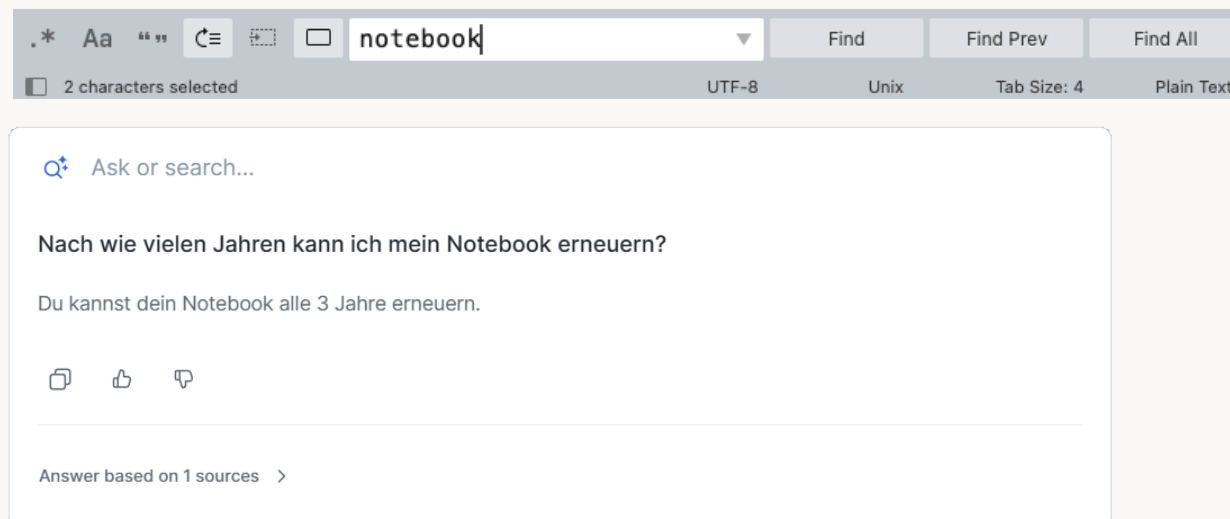


Recap: Not good enough?

- Tune text cleanup, segmentation, splitting
- HyDE or HyQE or alternative indexing
 - How many questions?
 - With or without summary
- Other approaches
 - Only generate summary
 - Extract “Intent” from user input and search by that
 - Transform document and query to a common search embedding
 - HyKSS: Hybrid Keyword and Semantic Search
<https://www.deg.byu.edu/papers/HyKSS.pdf>
- Always evaluate approaches with your own data & queries
- The actual / final approach is more involved
as it seems on the first glance

Conclusion

- Semantic search is a first and fast Generative AI business use-case
- Quality of results depend heavily on data quality and preparation pipeline
- RAG pattern can produce breathtaking good results without the need for user training



Thank you!



Demos:

<https://github.com/thinktecture-labs/basta-2024-advanced-rag>

Sebastian Gingter

<https://thinktecture.com/sebastian-gingter>

Real-World RAG:

Eigene Daten & Dokumente mit
semantischer Suche & LLMs erschließen

think
tecture

Slides & Code

<https://www.thinktecture.com/de/sebastian-gingter>

Sebastian Gingter

sebastian.gingter@thinktecture.com

Developer Consultant

